

Assignment-3

1. Recitation Exercises

Chapter-6

1. (a) Which of the three models with  $k$  predictors has the smallest training RSS?

The best subnet selection approach will have the smallest training RSS because it considers all the possible models. Other models mentioned use a greedy approach.

(b) Which of the three models with  $k$  predictors has the smallest test RSS?

We cannot determine if any of the three models have the small test RSS. The best subnet model may have low-training RSS since it considers all models. But this may overfit the model resulting in low test RSS.

(c) True or False

(i) True

(ii) True

(iii) False= sometimes it can be true but not always

(iv) False= sometimes it can be true but not always

(v) False= Best subnet selection may drop previously chosen predictors when a new one is added since it considers all possible subnets.

2. (a) The lasso, relative to least squares, is:

Option (iii) is correct. When it performs feature selection Lasso generates a sparse model that is less flexible.

(b) Repeat (a) for ridge regression relative to least squares.

Option (iii) is correct. When ridge performs shrinkage, it shrinks predictors that don't have a strong relationship with the target variable hence they are less flexible. It reduces the variance at the cost of an increase in bias.

(c) Repeat (a) for non-linear methods relative to least squares.

Option (ii) is correct. The non-linear model follows the observation more tightly than least squares hence, it's more flexible. It also reduces the bias at the cost of an increase in variance.

3. (a) As we increase  $s$  from 0, the training RSS will:

Option (iv) is correct. As we increase  $s$ , the model becomes more and more flexible as the restriction on  $\beta$  is reducing; hence, coefficients increase from 0 to their least square estimate values.

(b) Repeat (a) for test RSS

Option (ii) is correct. As the model is becoming more and more flexible the test RSS will reduce first and then start increasing when overfitting will start.

(c) Repeat (a) for variance.

Option (iii) is correct. Variance increases with the increase in model flexibility.

(d) Repeat (a) for (squared) bias.

Option (iv) is correct. Bias decreases with the increase in flexibility.

(e) Repeat (a) for the irreducible error.

Option (v) is correct. The irreducible error model does not depend on the value of  $s$ .

4. (a) As we increase  $\lambda$  from 0, the training RSS will:

Option (iii) is correct. As we increase  $\lambda$ , the model becomes less and less flexible because the restriction on  $\beta$  is increasing and hence coefficients value comes closer to 0.

(b) Repeat (a) for test RSS.

Option (ii) is correct. As the model becomes less flexible the test RSS will reduce first and then increase when overfitting starts.

(c) Repeat (a) for variance

Option (iv) is correct. Variance decreases with the decrease in model flexibility.

(d) Repeat (a) for (squared) bias.

Option (iii) is correct. Bias increases with a decrease in model flexibility.

(e) Repeat (a) for the irreducible error.

Option(v) is correct. The irreducible error model does not depend on  $\lambda$ .

5. (a) Write out the ridge regression optimization problem in this setting.

### Chapter-6

#### Recitation Exercises

5 a)  $n=2$   
 $p=2$   
 $x_{11} = x_{12}$   
 $y_1 + y_2 = 0$   
 $x_{12} + y_{22} = 0$   
 $x_{11} + x_{21} = 0$   
 $\beta_0 = 0$

#### Ridge Regression

$n=2$   
 $p=2$   
 $\beta_0 = 0$

a) 
$$\min \left[ (y_1 - \hat{\beta}_1 x_{11} - \hat{\beta}_2 x_{12})^2 + (y_2 - \hat{\beta}_1 x_{21} - \hat{\beta}_2 x_{22})^2 \right] + \lambda (\hat{\beta}_1^2 + \hat{\beta}_2^2)$$

(b) From a

$$(y_1^2 + \hat{\beta}_1^2 x_{11}^2 + \hat{\beta}_2^2 x_{12}^2 - 2\hat{\beta}_2 x_{12} y_1 - 2\hat{\beta}_1 x_{11} y_1 + 2\hat{\beta}_1 \hat{\beta}_2 x_{11} x_{12})$$

$$+ (y_2^2 + \hat{\beta}_1^2 x_{21}^2 + \hat{\beta}_2^2 x_{22}^2 - 2\hat{\beta}_2 x_{22} y_2 + 2\hat{\beta}_1 \hat{\beta}_2 x_{21} x_{22} - 2\hat{\beta}_1 x_{21} y_1)$$

$$+ \lambda (\hat{\beta}_1^2 + \hat{\beta}_2^2)$$

taking the partial derivative to  $\beta_1$  & setting eq to 0

$$\frac{\partial}{\partial \beta_1} : (2\hat{\beta}_1 x_{11}^2 - 2x_{11}y_1 + 2\hat{\beta}_2 x_{11}x_{12}) + (2\hat{\beta}_1 x_{21}^2 - 2x_{21}y_2 + 2\hat{\beta}_2 x_{21}x_{22}) + 2\tau\hat{\beta}_1 = 0$$

Let  $x_{11} = x_{12} = x_1$  and  $x_{21} = x_{22} = x_2$

$$(\hat{\beta}_1 x_1^2 - x_1 y_1 + \hat{\beta}_2 x_1^2) + (\hat{\beta}_1 x_2^2 - x_2 y_2 + \hat{\beta}_2 x_2^2) + \tau\hat{\beta}_1 = 0$$

$$\beta_1 (x_1^2 + x_2^2) + \beta_2 (x_1^2 + x_2^2) + \tau\beta_1 = x_1 y_1 + x_2 y_2$$

Adding  $2\hat{\beta}_1 x_1 x_2$  and  $2\hat{\beta}_2 x_1 x_2$  on both sides

$$\hat{\beta}_1 (x_1 + x_2)^2 + \hat{\beta}_2 (x_1 + x_2)^2 + \tau\hat{\beta}_1 = x_1 y_1 + x_2 y_2 + 2\hat{\beta}_1 x_1 x_2 + 2\hat{\beta}_2 x_1 x_2$$

Given  $x_1 + x_2 = 0$

$$\tau\hat{\beta}_1 = x_1 y_1 + x_2 y_2 + 2\hat{\beta}_1 x_1 x_2 + 2\hat{\beta}_2 x_1 x_2 \quad \text{--- (1)}$$

Taking partial derivative to  $\beta_2$  we get

$$\tau\hat{\beta}_2 = x_1 y_1 + x_2 y_2 + 2\hat{\beta}_1 x_1 x_2 + 2\hat{\beta}_2 x_1 x_2 \quad \text{--- (2)}$$

From (1) & (2) we get  $\tau\hat{\beta}_1 = \tau\hat{\beta}_2 \Rightarrow \beta_1 = \beta_2$

### c) Lasso optimization

$$\min \left[ (y_1 - \hat{\beta}_1 x_{11} - \hat{\beta}_2 x_{12})^2 + (y_2 - \hat{\beta}_1 x_{21} - \hat{\beta}_2 x_{22})^2 + \lambda (|\hat{\beta}_1| + |\hat{\beta}_2|) \right]$$

d) Replacing the penalty term from ridge regression the derivative term to  $\beta$

$$= -\frac{\partial}{\partial \beta} (\lambda |\beta|); \quad \frac{\partial}{\partial \beta} |\beta|$$

In ridge regression

$$\frac{\partial}{\partial \beta_1} \lambda |\beta_1| = \lambda \frac{\partial}{\partial \beta_1} |\beta_1|$$

$\beta_1$  and  $\beta_2$  can be both positive or both negative.

## 2. Recitation Exercises

### Chapter-7

2.(a) Suppose that a curve  $\hat{g}$  is computed to smoothly fit a set of  $n$  points using the following formula.

(a)  $\lambda = \infty, m = 0$ .

In this case,  $\hat{g}$  is 0 because due to the large smoothing parameter  $\hat{g}^0(x) \rightarrow 0$

(b)  $\lambda = \infty, m = 1$ .

In this case,  $\hat{g}$  is  $c$  because due to the large smoothing parameter  $\hat{g}^1(x) \rightarrow 0$

(c)  $\lambda = \infty, m = 2$ .

In this case,  $\hat{g}$  is  $cx + d$  because due to the large smoothing parameter  $\hat{g}^2(x) \rightarrow 0$

(d)  $\lambda = \infty, m = 3$ .

In this case,  $\hat{g}$  is  $cx^2 + dx + e$  because due to the large smoothing parameter  $\hat{g}^3(x) \rightarrow 0$

(e)  $\lambda = 0, m = 3$


In this case, the smoothing will have no effect as  $\lambda = 0$  and  $\hat{g}$  will interpolate the data set.

3.  $x = -2:2$

$$y = 1 + x + -2 * (x-1)^2 * I(x>1)$$

plot(x, y)

```
> x = -2:2
> y = 1 + x + -2 * (x-1)^2 * I(x>1)
> plot(x, y)
> |
```



4.  $x = -2:6$

$y = c(1 + 0 + 0, \# x = -2$

$1 + 0 + 0, \# x = -1$

$1 + 1 + 0, \# x = 0$

$1 + (1-0) + 0, \# x = 1$

$1 + (1-1) + 0, \# x = 2$

$1 + (0-0) + 0, \# x = 3$

$1 + (0) + 1, \# x = 4$


$1 + 0 + 1, \# x = 5$

$1 + 0 + 0 \# x = 6$

)

plot(x,y)

```
> x = -2:6
> y = c(1 + 0 + 0, # x = -2
+      1 + 0 + 0, # x = -1
+      1 + 1 + 0, # x = 0
+      1 + (1-0) + 0, # x = 1
+      1 + (1-1) + 0, # x = 2
+      1 + (0-0) + 0, # x = 3
+      1 + (0) + 1, # x = 4
+      1 + 0 + 1, # x = 5
+      1 + 0 + 0 # x = 6
+ )
> plot(x,y)
>
```



5. (a) As  $\lambda \rightarrow \infty$ , will  $\hat{g}_1$  or  $\hat{g}_2$  have the smaller training RSS?

$\hat{g}_2$  will have smaller training RSS as it will be a higher-order polynomial due to the order of penalty term.

(b) As  $\lambda \rightarrow \infty$ , will  $\hat{g}_1$  or  $\hat{g}_2$  have the smaller test RSS?

$\hat{g}_1$  will have a smaller test RSS value.

(c) For  $\lambda = 0$ , will  $\hat{g}_1$  or  $\hat{g}_2$  have the smaller training and test RSS?

Here  $\lambda = 0$  and we have  $\hat{g}_1 = \hat{g}_2$  so they will have equal training and test RSS.

## 2. Practicum Problems

### Problem#1

```
library(caret)
```

```
library(ggplot2)
```

```
data("mtcars")
```

```
colnames(mtcars)
```

```
> library(caret)
> library(ggplot2)
> data("mtcars")
> colnames(mtcars)
[1] "mpg" "cyl" "disp" "hp" "drat" "wt" "qsec" "vs" "am" "gear" "carb"
> |
```

### #80-20 training- testing split

```
SampleDataIndex = createDataPartition(y = mtcars$mpg, p = 0.8, list = FALSE)
```

```
TrainingDataMcars = mtcars[SampleDataIndex,]
```

```
TestingDtataMcars = mtcars[-SampleDataIndex,]
```

```
> #80-20 training- testing split
> SampleDataIndex = createDataPartition(y = mtcars$mpg, p = 0.8, list = FALSE)
> TrainingDataMcars = mtcars[SampleDataIndex,]
> TestingDtataMcars = mtcars[-SampleDataIndex,]
> |
```



## #fitting a linear model

```
LinearModelMcars1 = lm(mpg~., data = TrainingDataMcars)
```

```
summary(LinearModelMcars1)
```

```
> LinearModelMcars1 = lm(mpg~., data = TrainingDataMcars)
> summary(LinearModelMcars1)

Call:
lm(formula = mpg ~ ., data = TrainingDataMcars)

Residuals:
    Min       1Q   Median       3Q      Max
-2.6339 -1.6263 -0.1275  0.7253  4.6465

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  5.21282   20.13320   0.259  0.7988
cyl          0.57539    1.09241   0.527  0.6052
displ        0.01073    0.01813   0.592  0.5616
hp          -0.03350    0.02365  -1.417  0.1746
drat         1.51005    1.85994   0.812  0.4281
wt          -3.76657    1.93958  -1.942  0.0689
qsec         1.10211    0.75907   1.452  0.1647
vs          -0.87715    2.17482  -0.403  0.6917
am           3.09750    2.11809   1.462  0.1619

---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.622 on 17 degrees of freedom
Multiple R-squared:  0.88,    Adjusted R-squared:  0.8094
F-statistic: 12.47 on 10 and 17 DF,  p-value: 5.649e-06
```

## #coefficients values

```
LinearModelMcars1$coefficients
```

```
> #coefficients values
> LinearModelMcars1$coefficients
(Intercept)      cyl      displ      hp      drat      wt      qsec      vs      am      gear
5.21282197  0.57539004  0.01073177 -0.03349912  1.51005122 -3.76656990  1.10210634 -0.87714708  3.09749975  0.01575818
      carb
-0.01052182
> |
```

```
library(glmnet)
```

```
> library(glmnet)
> |
```

package 'RcppEigen' successfully unpacked and MD5 sums checked

The downloaded binary packages are in  
C:\Users\aaath\AppData\Local\Temp\RtmpoYy1WN\downloaded\_packages

[3/3] Installing glmnet...

Installing package into 'C:/Users/aaath/AppData/Local/R/win-library/4.2'  
(as 'lib' is unspecified)  
trying URL 'https://cran.rstudio.com/bin/windows/contrib/4.2/glmnet\_4.1-4.zip'  
Content type 'application/zip' length 2557243 bytes (2.4 MB)  
=====  
downloaded 2.4 MB

package 'glmnet' successfully unpacked and MD5 sums checked

The downloaded binary packages are in  
C:\Users\aaath\AppData\Local\Temp\RtmpoYy1WN\downloaded\_packages

✓ Package 'glmnet' successfully installed.

### #vector of 100 values

```
LambdaVector = 10^seq(5, -5, by = -.1)
```

```
> #vector of 100 values  
> LambdaVector = 10^seq(5, -5, by = -.1)  
> |
```

---

### #Extracting data for minimum value of lambda

```
y = TrainingDataMcars$mpg
```

```
x = model.matrix(mpg~., data = TrainingDataMcars)
```

```
> y = TrainingDataMcars$mpg  
> x = model.matrix(mpg~., data = TrainingDataMcars)  
> |
```

---

### #Use cross-validation (via cv.glmnet) to determine the minimum value for $\lambda$ - what do you obtain?

```
LambdaVectorCV = cv.glmnet(x, y, lambda = LambdaVector, alpha = 0 )
```

```
OptimalLambda = LambdaVectorCV$lambda.min
```

```
c(OptimalLambda)
```

```
> LambdaVectorCV = cv.glmnet(x, y, lambda = LambdaVector, alpha = 0 )  
warning message:  
Option grouped=FALSE enforced in cv.glmnet, since < 3 observations per fold  
> OptimalLambda = LambdaVectorCV$lambda.min  
> c(OptimalLambda)  
[1] 2.511886  
> |
```

---

Answer- Minimum value of lambda is 2.51186

### #Fitting the ridge regression


```
FitModel = glmnet(x, y, lambda = OptimalLambda, alpha = 0)
```

### #Plot training MSE as a function of $\lambda$ (you may also use $\log \lambda$ ).

```
plot(LambdaVectorCV)
```

```
> #Fitting the ridge regression  
> FitModel = glmnet(x, y, lambda = OptimalLambda, alpha = 0)  
> #Plot training MSE as a function of  $\lambda$  (you may also use  $\log \lambda$ ).  
> plot(LambdaVectorCV)  
> |
```





**#What is out-of-sample test set performance (using predict), and how do the coefficients differ versus the regular linear model?**

`coef(FitModel)`

```
> 'coef(FitModel)'
12 x 1 sparse Matrix of class "dgCMatrix"
      s0
(Intercept) 20.448198506
(Intercept) .
cyl         -0.287288332
disp        -0.005342594
hp          -0.013196073
drat         1.256065512
wt          -1.256736164
qsec         0.220717061
vs           0.305492479
am           1.685777455
gear         0.245686328
carb        -0.544637581
> |
```

**#Out of Sample test set performance for regular Linear model**

```
LinearModelMcarsPredict = predict(LinearModelMcars1, newdata = TestingDtataMcars, type = "response")
```

```
LinearModelMcars1Actual = TestingDtataMcars[, "mpg"]
```

```
MeanValueRegModel = mean((LinearModelMcars1Actual - LinearModelMcarsPredict)^2)
```

```
c(MeanValueRegModel)
```

```
> LinearModelMcarsPredict = predict(LinearModelMcars1, newdata = TestingDtataMcars, type = "response")
> LinearModelMcars1Actual = TestingDtataMcars[, "mpg"]
> MeanValueRegModel = mean((LinearModelMcars1Actual - LinearModelMcarsPredict)^2)
> c(MeanValueRegModel)
[1] 12.65157
> RegModel = model.matrix(mpg ~ 1, data = TestingDtataMcars)
```

**Answer-** Out of sample training MSE for Linear Model is 12.65157

**Has ridge regression performed shrinkage, variable selection, or both?**

**#Out of sample test set performance for regression Model**

```

RegModel = model.matrix(mpg~., data = TestingDtataMcars)

RegModelPredict = predict(FitModel, s = OptimalLambda, newx = RegModel)

RegModelActual = TestingDtataMcars[, "mpg"]

RegModelMean = mean((RegModelActual - RegModelPredict)^2)

print(RegModelMean)

> RegModel = model.matrix(mpg~., data = TestingDtataMcars)
> RegModelPredict = predict(FitModel, s = OptimalLambda, newx = RegModel)
> RegModelActual = TestingDtataMcars[, "mpg"]
> RegModelMean = mean((RegModelActual - RegModelPredict)^2)
> print(RegModelMean)
[1] 7.165049

```

**Answer-** Out of sample training MSE for Linear Model is 12.65157 and the out-of-sample Training MSE after ridge regression is 7.165049. So, it can be clearly seen that ridge regression has performed shrinkage.

## Problem#2

```

library(ggplot2)

library(caret)

data("swiss")

colnames(swiss)

[1] "Fertility" "Agriculture" "Examination" "Education" "Catholic" "Infant.Mortality"

```

## #80-20 training-testing split

```

SampleSwissIndex = createDataPartition(y = swiss$Fertility, p = 0.8, list = FALSE)

TrainingDataSwiss = swiss[SampleSwissIndex,]

TestingDataSwiss = swiss[-SampleSwissIndex,]

> #80-20 training-testing split
> SampleSwissIndex = createDataPartition(y = swiss$Fertility, p = 0.8, list = FALSE)
> TrainingDataSwiss = swiss[SampleSwissIndex,]
> TestingDataSwiss = swiss[-SampleSwissIndex,]
>

```

## #fitting a linear model

```

LinearModelSwiss1 = lm(Fertility~., data = TrainingDataSwiss)

summary(LinearModelSwiss1)

```

```
> LinearModelSwiss1 = lm(Fertility~., data = TrainingDataSwiss)
> summary(LinearModelSwiss1)
```

```
Call:
lm(formula = Fertility ~ ., data = TrainingDataSwiss)
```

```
Residuals:
    Min       1Q   Median       3Q      Max
-12.9988  -4.3723   0.2933   3.1141  13.9618
```

```
Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  62.04634   10.81152    5.739 2.08e-06 ***
Agriculture  -0.10866    0.07534   -1.442 0.158643
Examination  -0.18441    0.25427   -0.725 0.473397
Education    -0.75650    0.17782   -4.254 0.000163 ***
Catholic      0.11414    0.03475    3.284 0.002424 **
Infant.Mortality 1.03405    0.36984    2.796 0.008562 **
```

```
---
Catholic      0.11414    0.03475    3.284 0.002424 **
Infant.Mortality 1.03405    0.36984    2.796 0.008562 **
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 6.675 on 33 degrees of freedom
Multiple R-squared:  0.7477,    Adjusted R-squared:  0.7095
F-statistic: 19.56 on 5 and 33 DF,  p-value: 5.072e-09
```

```
> |
```

#### #coefficient values

#### #What are the associated coefficient values for relevant features?

```
LinearModelSwiss1$coefficients
```

```
> LinearModelSwiss1$coefficients
(Intercept)  Agriculture  Examination  Education  Catholic Infant.Mortality
 62.0463423   -0.1086631   -0.1844132   -0.7564982   0.1141418   1.0340487
```

```
library(glmnet)
```

#### #vector of 100 values

```
LambdaVectorSwiss = 10^seq(5, -5, by = -.1)
```

#### #Extracting x and y values from training data

```
y = TrainingDataSwiss$Fertility
```

```
x = model.matrix(Fertility~., data = TrainingDataSwiss)
```

```
> #vector of 100 values
> LambdaVectorSwiss = 10^seq(5, -5, by = -.1)
> #Extracting x and y values from training data
> y = TrainingDataSwiss$Fertility
> x = model.matrix(Fertility~., data = TrainingDataSwiss)
> |
```

#### #Use cross-validation (via cv.glmnet) to determine the minimum value for $\lambda$ - what do you obtain?

```
LambdaVectorSwissCV = cv.glmnet(x, y, lambda = LambdaVectorSwiss, alpha = 1)
```

```
OptimalLambdaSwiss = LambdaVectorSwissCV$lambda.min
```

c(OptimalLambdaSwiss)

```
> LambdaVectorSwissCV = cv.glmnet(x, y, lambda = LambdaVectorSwiss, alpha = 1)
> OptimalLambdaSwiss = LambdaVectorSwissCV$lambda.min
> c(OptimalLambdaSwiss)
[1] 0.6309573
>
```

**Answer-** Minimum value of lambda is 0.6309573


## #fitting lasso regression

```
FitModelSwiss = glmnet(x, y, lambda = OptimalLambdaSwiss, alpha = 1)
```

**#Plot training MSE as a function of  $\lambda$  (you may also use  $\log \lambda$ ).**

```
plot(LambdaVectorSwissCV)
```

```
> #fitting lasso regression
> FitModelSwiss = glmnet(x, y, lambda = optimalLambdaswiss, alpha = 1)
> #Plot training MSE as a function of  $\lambda$  (you may also use  $\log \lambda$ ).
> plot(LambdaVectorSwissCV)
>
```



**#What is out-of-sample test set performance (using predict), and how do the coefficients differ versus the regular linear model?**

```
coef(FitModelSwiss)
```

```
> plot(LambdaVectorSwissV)
> coef(FitModelSwiss)
7 x 1 sparse matrix of class "dgCMatrix"
      s0
(Intercept) 54.49765837
(Intercept) .
Agriculture .
Examination -0.06718265
Education -0.60898884
Catholic 0.10008801
Infant.Mortality 0.97759315
>
```

### #Out of Sample test set performance for regular Linear model

```
LinearModelSwissPredict = predict(LinearModelSwiss1, newdata = TestingDataSwiss, type = "response")  
LinearModelSwissActual = TestingDataSwiss[, "Fertility"]  
LinearModelSwissMean = mean((LinearModelSwissActual - LinearModelSwissPredict)^2)  
print(LinearModelSwissMean)
```

```
> LinearModelSwissPredict = predict(LinearModelSwiss1, newdata = TestingDataSwiss, type = "response")  
> LinearModelSwissActual = TestingDataSwiss[, "Fertility"]  
> LinearModelSwissMean = mean((LinearModelSwissActual - LinearModelSwissPredict)^2)  
> print(LinearModelSwissMean)  
[1] 89.6039  
>
```

Out of sample training MSE for regular Linear Model = 89.6039

### #Has lasso regression performed shrinkage, variable selection, or both?

### #Out of sample test set performance for lasso regression model

```
LassoRegModel = model.matrix(Fertility~., data = TestingDataSwiss)  
LassoRegModelPredict = predict(FitModelSwiss, s = OptimalLambdaSwiss, newx = LassoRegModel)  
LassoRegModelActual = TestingDataSwiss[, "Fertility"]  
LassoRegModelMean = mean((LassoRegModelActual - LassoRegModelPredict)^2)  
print(LassoRegModelMean)
```

```
> LassoRegModel = model.matrix(Fertility~., data = TestingDataSwiss)  
> LassoRegModelPredict = predict(FitModelSwiss, s = OptimalLambdaSwiss, newx = LassoRegModel)  
> LassoRegModelActual = TestingDataSwiss[, "Fertility"]  
> LassoRegModelMean = mean((LassoRegModelActual - LassoRegModelPredict)^2)  
> print(LassoRegModelMean)  
[1] 114.9967  
>
```

**Answer-** After performing Lasso regression the out-of-sample MSE has raised from 89.6039 to 114.9967. Lasso regression usually performs variable selection but, in this case, it performs shrinkage as the coefficients have shrunk to some extent.

### Problem#3

```
library(readxl)  
library(magrittr)  
ConcreteData1 = read_excel("C:\\Users\\aasth\\Downloads\\Concrete_Data.xls")  
summary(ConcreteData1)
```

```

> library(readxl)
> library(magrittr)
> ConcreteData1 = read_excel("c:\\Users\\aasth\\Downloads\\Concrete_Data.xls")
> summary(ConcreteData1)
Cement (component 1)(kg in a m^3 mixture) Blast Furnace Slag (component 2)(kg in a m^3 mixture)
Min. :102.0 Min. : 0.0
1st Qu.:192.4 1st Qu.: 0.0
Median :272.9 Median : 22.0
Mean :281.2 Mean : 73.9
3rd Qu.:350.0 3rd Qu.:142.9
Max. :540.0 Max. :359.4
Fly Ash (component 3)(kg in a m^3 mixture) Water (component 4)(kg in a m^3 mixture)
Min. : 0.00 Min. :121.8
1st Qu.: 0.00 1st Qu.:164.9
Median : 0.00 Median :185.0
Mean : 54.19 Mean :181.6
3rd Qu.:118.27 3rd Qu.:192.0
Max. :200.10 Max. :247.0
Superplasticizer (component 5)(kg in a m^3 mixture) Coarse Aggregate (component 6)(kg in a m^3 mixture)
Max. :540.0 Max. :359.4
Fly Ash (component 3)(kg in a m^3 mixture) Water (component 4)(kg in a m^3 mixture)
Min. : 0.00 Min. :121.8
1st Qu.: 0.00 1st Qu.:164.9
Median : 0.00 Median :185.0
Mean : 54.19 Mean :181.6
3rd Qu.:118.27 3rd Qu.:192.0
Max. :200.10 Max. :247.0
Superplasticizer (component 5)(kg in a m^3 mixture) Coarse Aggregate (component 6)(kg in a m^3 mixture)
Min. : 0.000 Min. : 801.0
1st Qu.: 0.000 1st Qu.: 932.0
Median : 6.350 Median : 968.0
Mean : 6.203 Mean : 972.9
3rd Qu.:10.160 3rd Qu.:1029.4
Max. :32.200 Max. :1145.0
Fine Aggregate (component 7)(kg in a m^3 mixture) Age (day) Concrete compressive strength(MPa, megapascals)
Min. :594.0 Min. : 1.00 Min. : 2.332
1st Qu.:731.0 1st Qu.: 7.00 1st Qu.:23.707
Median :779.5 Median :28.00 Median :34.443
Mean :773.6 Mean : 45.66 Mean :35.818
3rd Qu.:824.0 3rd Qu.: 56.00 3rd Qu.:46.136
Max. :992.6 Max. :365.00 Max. :82.599

```

colnames(ConcreteData1)

```

> colnames(ConcreteData1)
[1] "Cement (component 1)(kg in a m^3 mixture)" "Blast Furnace slag (component 2)(kg in a m^3 mixture)"
[3] "Fly Ash (component 3)(kg in a m^3 mixture)" "Water (component 4)(kg in a m^3 mixture)"
[5] "Superplasticizer (component 5)(kg in a m^3 mixture)" "Coarse Aggregate (component 6)(kg in a m^3 mixture)"
[7] "Fine Aggregate (component 7)(kg in a m^3 mixture)" "Age (day)"
[9] "Concrete compressive strength(MPa, megapascals)"

```

library(mgcv)

library(stringr)

names(ConcreteData1)[1] = "cement" #c1

names(ConcreteData1)[2] = "slag" #c2

names(ConcreteData1)[3] = "ash" #c3

names(ConcreteData1)[4] = "water" #c4

names(ConcreteData1)[5] = "superplasticizer" #c5

names(ConcreteData1)[6] = "coarse" #c6

names(ConcreteData1)[7] = "fine"

names(ConcreteData1)[8] = "age"

names(ConcreteData1)[9] = "ccs"

names(ConcreteData1)



```

[5] Concrete compressive strength (MPa, megapascals)
> library(mgcv)
> library(stringr)
> names(ConcreteData1)[1] = "cement"      #c1
> names(ConcreteData1)[2] = "slag"         #c2
> names(ConcreteData1)[3] = "ash"          #c3
> names(ConcreteData1)[4] = "water"        #c4
> names(ConcreteData1)[5] = "superplasticizer" #c5
> names(ConcreteData1)[6] = "coarse"       #c6
> names(ConcreteData1)[7] = "fine"
> names(ConcreteData1)[8] = "age"
> names(ConcreteData1)[9] = "ccs"
> names(ConcreteData1)
[1] "cement"      "slag"         "ash"          "water"        "superplasticizer" "coarse"
[7] "fine"        "age"          "ccs"
> |

```

### #creating a generalized additive model as linear model for C1-C6

```
LinearModelAddict = gam(ccs ~ cement + slag + ash + water + superplasticizer + coarse, data = ConcreteData1 )
```

```
summary(LinearModelAddict)
```

```

> #creating a generalized additive model as linear model for C1-C6
> LinearModelAddict = gam(ccs ~ cement + slag + ash + water + superplasticizer + coarse, data = ConcreteData1 )
> summary(LinearModelAddict)

```

```

Family: gaussian
Link function: identity

```

```

Formula:
ccs ~ cement + slag + ash + water + superplasticizer + coarse

```

```

Parametric coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)   5.326997  10.510518   0.507 0.612387
cement        0.108256   0.005214  20.761 < 2e-16 ***
slag          0.079357   0.006193  12.814 < 2e-16 ***
ash           0.055928   0.009287   6.022 2.4e-09 ***
water        -0.103871   0.027796  -3.737 0.000197 ***
superplasticizer 0.356016   0.110251   3.229 0.001281 **
coarse        0.008027   0.006272   1.280 0.200940
---

```

```

coarse          0.008027   0.006272   1.280 0.200940
---

```

```

signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

```

R-sq.(adj) =  0.445   Deviance explained = 44.9%
GCV = 155.83   Scale est. = 154.77    n = 1030

```

**Answer-** The value of R2 here is 0.445

### #creating generalized additive model as linear model for all values

```
LinearModelAddict1 <- gam(ccs ~ cement + slag + ash + water + superplasticizer + coarse + fine +
```

```
age, data = ConcreteData1)
```

```
summary(LinearModelAddict1)
```

```

> #creating generalized additive model as linear model for all values
> LinearModelAddict1 <- gam(ccs ~ cement + slag + ash + water + superplasticizer + coarse + fine +
+ age, data = ConcreteData1)
> summary(LinearModelAddict1)

```

```

Family: gaussian
Link function: identity

```

```

Formula:
ccs ~ cement + slag + ash + water + superplasticizer + coarse +
      fine + age

```

```

Parametric coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  -23.163756  26.588421  -0.871 0.383851
cement        0.119785   0.008489  14.110 < 2e-16 ***

```

```

Parametric coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  -23.163756  26.588421  -0.871 0.383851
cement        0.119785   0.008489  14.110 < 2e-16 ***
slag          0.103847   0.010136  10.245 < 2e-16 ***
ash           0.087943   0.012585   6.988 5.03e-12 ***
water        -0.150298   0.040179  -3.741 0.000194 ***
superplasticizer 0.290687   0.093460   3.110 0.001921 **
coarse        0.018030   0.009394   1.919 0.055227 .
fine          0.020154   0.010703   1.883 0.059968 .
age           0.114226   0.005427  21.046 < 2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

R-sq.(adj) =  0.612   Deviance explained = 61.5%
GCV = 109.11   Scale est. = 108.16     n = 1030
> |

```

**Answer-** In this the R2 value for linear model is 0.612 respectively.

### #creating a generalized additive model as non linear for C1-C6

```

NonLinearModelAddict = gam(ccs ~ s(cement)+ s(slag)+ s(water)+ s(ash)+
s(superplasticizer)+ s(coarse), data = ConcreteData1)

```

```
summary(NonLinearModelAddict)
```

```

GCV = 109.11   Scale est. = 108.16     n = 1030
> NonLinearModelAddict = gam(ccs ~ s(cement)+ s(slag)+ s(water)+ s(ash)+
+ s(superplasticizer)+ s(coarse), data = ConcreteData1)
> summary(NonLinearModelAddict)

```

```

Family: gaussian
Link function: identity

```

```

Formula:
ccs ~ s(cement) + s(slag) + s(water) + s(ash) + s(superplasticizer) +
      s(coarse)

```

```

Parametric coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  35.8178     0.3566  100.4 <2e-16 ***
---

```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Approximate significance of smooth terms:
```

```

Parametric coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  35.8178     0.3566  100.4 <2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

```
Approximate significance of smooth terms:
```

|                     | edf   | Ref.df | F      | p-value      |
|---------------------|-------|--------|--------|--------------|
| s(cement)           | 4.464 | 5.513  | 69.530 | < 2e-16 ***  |
| s(slag)             | 2.088 | 2.578  | 48.091 | < 2e-16 ***  |
| s(water)            | 8.567 | 8.936  | 13.504 | < 2e-16 ***  |
| s(ash)              | 5.332 | 6.404  | 1.784  | 0.101        |
| s(superplasticizer) | 7.133 | 8.143  | 5.498  | 1.22e-06 *** |
| s(coarse)           | 1.000 | 1.000  | 0.018  | 0.892        |

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```

R-sq.(adj) =  0.531   Deviance explained = 54.4%
GCV = 134.84   Scale est. = 130.96     n = 1030
> |

```

**Answer-** R2 value for non-linear model is 0.531

### #Creating a generalised additive model as non- linear for all values

```
NonLinearModelAddict2 <- gam(ccs ~ s(cement) + s(slag) + s(ash) + s(water) + s(superplasticizer) +
s(coarse) + s(fine) + s(age), data = ConcreteData1)

summary(NonLinearModelAddict2)
```

```
> #Creating a generalised additive model as non- linear for all values
> NonLinearModelAddict2 <- gam(ccs ~ s(cement) + s(slag) + s(ash) + s(water) + s(superplasticizer) +
+ s(coarse) + s(fine) + s(age), data = ConcreteData1)
> summary(NonLinearModelAddict2)
```

```
Family: gaussian
Link function: identity
```

```
Formula:
ccs ~ s(cement) + s(slag) + s(ash) + s(water) + s(superplasticizer) +
      s(coarse) + s(fine) + s(age)
```

```
Parametric coefficients:
```

```
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  35.8178     0.1675    213.9  <2e-16 ***
---
```

```
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  35.8178     0.1675    213.9  <2e-16 ***
---
```

```
signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Approximate significance of smooth terms:
```

|                     | edf   | Ref.df | F       | p-value  |     |
|---------------------|-------|--------|---------|----------|-----|
| s(cement)           | 8.223 | 8.830  | 48.690  | < 2e-16  | *** |
| s(slag)             | 8.114 | 8.757  | 25.041  | < 2e-16  | *** |
| s(ash)              | 8.256 | 8.817  | 9.354   | < 2e-16  | *** |
| s(water)            | 8.742 | 8.973  | 26.116  | < 2e-16  | *** |
| s(superplasticizer) | 8.039 | 8.743  | 10.845  | < 2e-16  | *** |
| s(coarse)           | 7.904 | 8.673  | 3.403   | 0.000593 | *** |
| s(fine)             | 8.618 | 8.951  | 18.435  | < 2e-16  | *** |
| s(age)              | 8.559 | 8.900  | 365.119 | < 2e-16  | *** |

```
signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
R-sq.(adj) =  0.896   Deviance explained = 90.3%
```

```
GCV = 30.914   Scale est. = 28.89       n = 1030
```

```
> |
```

Answer- In this R2 is 0.896 respectively.

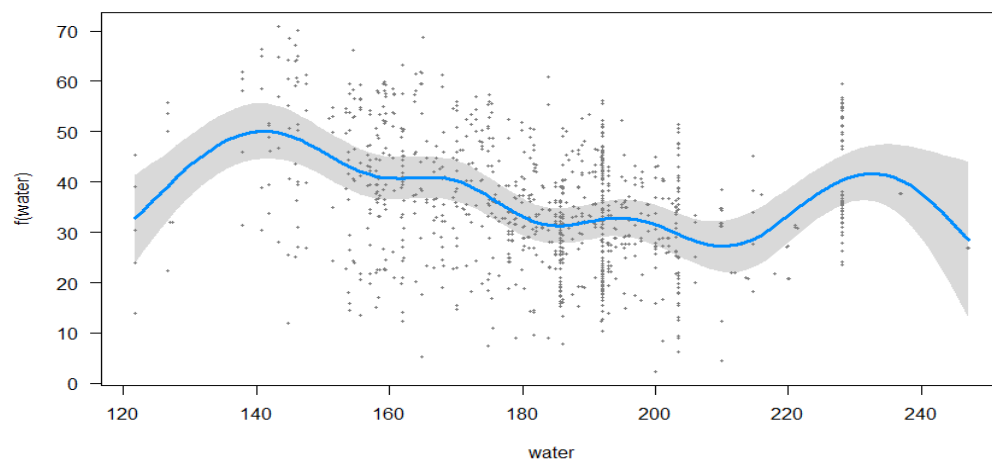


**# Visualize the regression using the visreg package, showing the fit as a function of each predictor with confidence intervals**




```
library(visreg)
```

```
visreg(NonLinearModelAddict)
```

```
visreg(NonLinearModelAddict2)
```



```
> library(visreg)
> visreg(NonLinearModelAddict)
Hit <Return> to see next plot:
Hit <Return> to see next plot:
Hit <Return> to see next plot:
Hit <Return> to see next plot:
Hit <Return> to see next plot:
>
> |
```











```
visreg(NonLinearModelAddict2)
```

```
>  
> visreg(NonLinearModelAddict2)  
Hit <Return> to see next plot:  
Hit <Return> to see next plot:  
Hit <Return> to see next plot:  
Hit <Return> to see next plot:  
Hit <Return> to see next plot:  
Hit <Return> to see next plot:  
Hit <Return> to see next plot:  
>
```









**Submitted By: -**

**Aastha Dhir**

**CWID-A20468022**

**adhir2@hawk.iit.edu**